

MySQL - HackerRank Notebook

Pham Dang Hoang Thien

February 8, 2026

Student, Ho Chi Minh City, Vietnam
e-mail: hoangthien312006@gmail.com

Contents

1	Introduction to MySQL	2
2	Basic SQL	2
2.1	Triangle Classification Query	2
2.2	Revising Aggregations	4
2.2.1	The Count Function	4
2.2.2	The Sum Function	4
2.2.3	Averages	5
2.2.4	Average Population	5
2.2.5	Japan Population	6
2.2.6	Population Density Difference	6
2.3	The Blunder	7
2.4	Top Earners	9
2.5	Weather Observation	11
2.5.1	Station 2	11
2.5.2	Station 13	12
2.5.3	Station 14	12
2.5.4	Station 15	13
2.5.5	Station 16	13
2.5.6	Station 17	13
2.6	Population Census	14
2.7	African Cities	15
2.8	Average Population of Each Continent	15
2.9	Draw The Triangle 1	16
2.10	Draw The Triangle 2	17
3	Intermediate SQL	18
3.1	The PADS	18
3.2	Occupations	21
3.3	Binary Tree Nodes	22
3.4	New Companies	24

1 Introduction to MySQL

2 Basic SQL

2.1 Triangle Classification Query

Write a query identifying the type of each record in the `TRIANGLES` table using its three side lengths. Output one of the following statements for each record in the table:

- **Equilateral:** It's a triangle with 3 sides of equal length.
- **Isosceles:** It's a triangle with 2 sides of equal length.
- **Scalene:** It's a triangle with 3 sides of differing lengths.
- **Not A Triangle:** The given values of A, B, and C don't form a triangle.

Input Format:

<i>Column</i>	<i>Type</i>
<i>A</i>	<i>Integer</i>
<i>B</i>	<i>Integer</i>
<i>C</i>	<i>Integer</i>

The `TRIANGLES` table contains three columns: A, B, and C, representing the lengths of the sides of a triangle.

Sample Input:

A	B	C
20	20	23
20	20	20
20	21	22
13	14	30

Sample Output:

Isosceles
Equilateral
Scalene
Not A Triangle

Solution Hint

Use a **CASE** expression to classify each triangle in this order:

1. Check **Not A Triangle** first using the triangle inequality:

$$A + B \leq C \text{ or } A + C \leq B \text{ or } B + C \leq A$$

2. If valid, check **Equilateral**: $A = B \text{ AND } B = C$
3. Check **Isosceles**: exactly two sides are equal
4. Otherwise, it is **Scalene**

Syntax elements used:

- CASE WHEN ... THEN ... ELSE ... END
- Logical operators: AND, OR
- Comparison operators: =, <=

SQL Query

```
1 SELECT
2   CASE
3     WHEN A + B <= C OR A + C <= B OR B + C <= A THEN 'Not A
4       Triangle'
5     WHEN A = B AND B = C THEN 'Equilateral'
6     WHEN A = B OR B = C OR A = C THEN 'Isosceles'
7     ELSE 'Scalene'
8   END AS triangle_type
FROM TRIANGLES;
```

Why this condition order?

- Put **Not A Triangle** first to filter out invalid side combinations.
- Then classify only valid triangles as equilateral, isosceles, or scalene.

2.2 Revising Aggregations

Input Format: The CITY table is described as follows:

<i>Fields</i>	<i>Type</i>
<i>ID</i>	<i>Number</i>
<i>NAME</i>	<i>VARCHAR2(17)</i>
<i>COUNTRYCODE</i>	<i>VARCHAR2(3)</i>
<i>DISTRICT</i>	<i>VARCHAR2(20)</i>
<i>POPULATION</i>	<i>Number</i>

2.2.1 The Count Function

Query a count of the number of cities in CITY having a Population larger than 100,000.

Solution Hint

The query should count the number of rows in the CITY table where the POPULATION is greater than 100,000.

COUNT(*) is used to count all rows that meet the specified condition in the WHERE clause. AS **city_count** renames the output column for clarity.

SQL Query

```
1 SELECT COUNT(*) AS city_count
2 FROM CITY
3 WHERE POPULATION > 100000;
```

2.2.2 The Sum Function

Query the total population of all cities in CITY where District is California.

Solution Hint

The query should sum the POPULATION column in the CITY table where the DISTRICT is 'California'.

SUM(POPULATION) is used to calculate the total population of rows that meet the specified condition in the WHERE clause. AS **total_population** renames the output column for clarity.

SQL Query

```
1 SELECT SUM(POPULATION) AS total_population
2 FROM CITY
3 WHERE DISTRICT = 'California';
```

2.2.3 Averages

Query the average population of all cities in CITY where District is California.

Solution Hint

The query should calculate the average of the POPULATION column in the CITY table where the DISTRICT is 'California'.

AVG(POPULATION) is used to calculate the average population of rows that meet the specified condition in the WHERE clause. AS *average_population* renames the output column for clarity.

SQL Query

```
1 SELECT AVG(POPULATION) AS average_population
2 FROM CITY
3 WHERE DISTRICT = 'California';
```

2.2.4 Average Population

Query the average population for all cities in CITY, rounded down to the nearest integer.

Solution Hint

The query should calculate the average of the POPULATION column in the CITY table and round it down to the nearest integer.

- Use **AVG(POPULATION)** to get the average population.
- Use **ROUND(...)** to round the average population to the nearest integer.
- AS *average_population* renames the output column for clarity.

SQL Query

```
1 SELECT ROUND(AVG(POPULATION)) AS average_population
2 FROM CITY;
```

2.2.5 Japan Population

Query the total population of cities in CITY where CountryCode is JPN.

Solution Hint

The query should sum the POPULATION column in the CITY table where the COUNTRYCODE is 'JPN'.

SUM(POPULATION) is used to calculate the total population of rows that meet the specified condition in the WHERE clause. AS ***total_population*** renames the output column for clarity.

SQL Query

```
1 SELECT SUM(POPULATION) AS total_population
2 FROM CITY
3 WHERE COUNTRYCODE = 'JPN';
```

2.2.6 Population Density Difference

Query the difference between the maximum and minimum populations in CITY.

Solution Hint

The query should calculate the difference between the maximum and minimum values of the POPULATION column in the CITY table.

- Use ***MAX(POPULATION)*** to get the maximum population.
- Use ***MIN(POPULATION)*** to get the minimum population.
- AS ***pop_diff*** renames the output column for clarity.

SQL Query

```
1 SELECT MAX(POPULATION) – MIN(POPULATION) AS pop_diff
2 FROM CITY;
```

2.3 The Blunder

Samantha was tasked with calculating the average monthly salaries for all employees in the EMPLOYEES table, but did not realize her keyboard's 0 key was broken until after completing the calculation. She wants your help finding the difference between her miscalculation (using salaries with any zeros removed), and the actual average salary. Write a query calculating the amount of error (i.e.: *actual-miscalculated* average monthly salaries), and round it up to the next integer. **Input Format**

The EMPLOYEES table is described as follows:

<i>Column</i>	<i>Type</i>
<i>ID</i>	<i>Integer</i>
<i>Name</i>	<i>String</i>
<i>Salary</i>	<i>Integer</i>

Note: Salary is per month.

Constraints

$$1000 < \text{Salary} < 10^5$$

Sample Input

<i>ID</i>	<i>Name</i>	<i>Salary</i>
1	Kristeen	1420
2	Ashley	2006
3	Julia	2210
4	Maria	3000

Sample Output

2061

Explanation

The table below shows the salaries without zeros as they were entered by Samantha:

<i>ID</i>	<i>Name</i>	<i>Salary</i>
1	Kristeen	142
2	Ashley	26
3	Julia	221
4	Maria	3

Samantha computes an average salary of 98.00. The actual average salary is 2159.00.

The resulting error between the two calculations is

$$2159.00 - 98.00 = 2061.00$$

Since it is equal to the integer 2061, it does not get rounded up.

Solution Hint

The query should calculate the difference between the actual average salary and the miscalculated average salary (with zeros removed). The difference should be rounded up to the next integer.

- ***AVG(Salary)*** is used to calculate the actual average salary.
- ***REPLACE(CAST(Salary AS CHAR), '0', '')*** is used to remove all zeros from the salary by converting it to a string, replacing '0' with an empty string, and then casting it back to an unsigned integer. /để hiểu là bỏ hết số 0 trong salary (mô phỏng bàn phím hỏng) đi rồi tính trung bình.
- ***AS UNSIGNED*** ensures that the modified salary is a positive integer for the average calculation. /nhận giá trị từ 0 trở lên.

SQL Query

```
1 SELECT CEIL(  
2     AVG(Salary) - AVG(CAST(REPLACE(CAST(Salary AS CHAR),  
3         '0', '') AS UNSIGNED))  
4 ) AS error  
5 FROM EMPLOYEES;
```

2.4 Top Earners

We define an employee's total earnings to be their monthly *salary* $\times$ *months* worked, and the maximum total earnings to be the maximum total earnings for any employee in the Employee table. Write a query to find the maximum total earnings for all employees as well as the total number of employees who have maximum total earnings. Then print these values as 2 space-separated integers. **Input Format:** The **EMPLOYEES** table containing employee data for a company is described as follows:

<i>Column</i>	<i>Type</i>
<i>employee_id</i>	<i>Integer</i>
<i>Name</i>	<i>String</i>
<i>Months</i>	<i>Integer</i>
<i>Salary</i>	<i>Integer</i>

where *employee_id* is an employee's ID number, *name* is their name, *months* is the total number of months they've been working for the company, and *salary* is their monthly salary.

Sample Input:

employee_id	name	months	salary
12228	Rose	15	1968
33645	Angela	1	3443
45692	Frank	17	1608
56118	Patrick	7	1345
59725	Lisa	11	2330
74197	Kimberly	16	4372
78454	Bonnie	8	1771
83565	Michael	6	2017
98607	Todd	5	3396
99989	Joe	9	3573

Sample Output

69952 1

Explanation

The table and earnings data are depicted below:

employee_id	name	months	salary	earnings
12228	Rose	15	1968	29520
33645	Angela	1	3443	3443
45692	Frank	17	1608	27336
56118	Patrick	7	1345	9415
59725	Lisa	11	2330	25630
74197	Kimberly	16	4372	69952
78454	Bonnie	8	1771	14168
83565	Michael	6	2017	12102
98607	Todd	5	3396	16980
99989	Joe	9	3573	32157

The maximum earnings value is 69952. The only employee with earnings = 69952 is Kimberly, so we print the maximum earnings value (69952) and the count of employees who earned \$69952 (1) as two space-separated values.

Solution Hint

eg. if total earnings is: 12000, 15000, 15000, 8000. Then the max total earnings is 15000, and there are 2 employees with that earning. So the output is "15000 2".

- 1. Calculate total earnings for each employee as *months * salary*.
- 2. Find the maximum total earnings using *MAX(months * salary)*.
- 3. Count the number of employees with total earnings equal to the maximum using *COUNT(*)* with a WHERE clause.

SQL Query

```
1 SELECT
2     MAX(months * salary) AS max_earnings,
3     COUNT(*) AS employee_count
4 FROM EMPLOYEES
5 WHERE (months * salary) = (
```

```

6  SELECT MAX(months * salary)
7  FROM EMPLOYEES
8  );

```

2.5 Weather Observation

Input Format: The STATION table is described as follows:

<i>Column</i>	<i>Type</i>
<i>ID</i>	<i>Number</i>
<i>CITY</i>	<i>VARCHAR2(21)</i>
<i>STATE</i>	<i>VARCHAR2(2)</i>
<i>LAT_N</i>	<i>Number</i>
<i>LONG_W</i>	<i>Number</i>

where LAT_N is the northern latitude and LONG_W is the western longitude.

2.5.1 Station 2

Query the following two values from STATION:

1. The sum of all values in LAT_N rounded to scale of 2 decimal places.
2. The sum of all values in LONG_W rounded to scale of 2 decimal places.

Output Format: Your results must be in the form:

lat lon

where *lat* is the sum of all values in LAT_N and *lon* is the sum of all values in LONG_W. Both results must be rounded to a scale of 2 decimal places.

Solution Hint

- Use **SUM(LAT_N)** to calculate the sum of all values in LAT_N.
- Use **SUM(LONG_W)** to calculate the sum of all values in LONG_W.
- Use **ROUND(..., 2)** to round both sums to a scale of 2 decimal places.
- AS **lat** and AS **lon** rename the output columns for clarity.

SQL Query

```
1 SELECT
2     ROUND(SUM(LAT_N), 2) AS lat,
3     ROUND(SUM(LONG_W), 2) AS lon
4 FROM STATION;
```

2.5.2 Station 13

Query the sum of Northern Latitudes (LAT_N) from STATION having values greater than 38.7880 and less than 137.2345. Truncate your answer to 4 decimal places.

SQL Query

```
1 SELECT ROUND(SUM(LAT_N), 4) AS total_latitude
2 FROM STATION
3 WHERE LAT_N > 38.7880 AND LAT_N < 137.2345;
```

2.5.3 Station 14

Query the greatest value of the Northern Latitudes (LAT_N) from STATION that is less than 137.2345. Truncate your answer to 4 decimal places.

Solution Hint

With x as a number and n as the number of decimal places:

- ***ROUND***(x , n) is rounding function that rounds the number x to n decimal places. eg. ***ROUND***(3.14159, 4) = 3.1416
- ***TRUNCATE***(x , n) is truncating function that truncates the number x to n decimal places. eg. ***TRUNCATE***(3.14159, 4) = 3.1415

SQL Query

```
1 SELECT TRUNCATE(MAX(LAT_N), 4) AS max_latitude
2 FROM STATION
3 WHERE LAT_N < 137.2345;
```

2.5.4 Station 15

Query the Western Longitude (LONG_W) for the largest Northern Latitude (LAT_N) in STATION that is less than 137.2345. Round your answer to 4 decimal places.

SQL Query

```
1 SELECT TRUNCATE(LONG_W, 4)
2 FROM STATION
3 WHERE LAT_N = (
4     SELECT MAX(LAT_N)
5     FROM STATION
6     WHERE LAT_N < 137.2345
7 );
```

2.5.5 Station 16

Query the smallest Northern Latitude (LAT_N) from STATION that is greater than 38.7780. Round your answer to 4 decimal places.

SQL Query

```
1 SELECT ROUND(MIN(LAT_N), 4)
2 FROM STATION
3 WHERE LAT_N > 38.7780;
```

2.5.6 Station 17

Query the Western Longitude (LONG_W) where the smallest Northern Latitude (LAT_N) in STATION is greater than 38.7780. Round your answer to 4 decimal places.

SQL Query

```
1 SELECT ROUND(LONG_W, 4)
2 FROM STATION
3 WHERE LAT_N = (
4     SELECT MIN(LAT_N)
5     FROM STATION
```

```

6 WHERE LAT_N > 38.7780
7 );

```

2.6 Population Census

Given the CITY and COUNTRY tables, query the sum of the populations of all cities where the CONTINENT is 'Asia'.

Note: CITY.CountryCode and COUNTRY.Code are matching key columns.

Input Format:

<i>Column</i>	<i>Type</i>
<i>ID</i>	<i>NUMBER</i>
<i>NAME</i>	<i>VARCHAR2(17)</i>
<i>COUNTRYCODE</i>	<i>VARCHAR2(3)</i>
<i>DISTRICT</i>	<i>VARCHAR2(20)</i>
<i>POPULATION</i>	<i>NUMBER</i>

<i>Column</i>	<i>Type</i>
<i>CODE</i>	<i>VARCHAR2(3)</i>
<i>NAME</i>	<i>VARCHAR2(44)</i>
<i>CONTINENT</i>	<i>VARCHAR2(13)</i>
<i>REGION</i>	<i>VARCHAR2(25)</i>
<i>SURFACEAREA</i>	<i>NUMBER</i>
<i>INDEPYEAR</i>	<i>VARCHAR2(5)</i>
<i>POPULATION</i>	<i>NUMBER</i>
<i>LIFEEXPECTANCY</i>	<i>VARCHAR2(4)</i>
<i>GNP</i>	<i>NUMBER</i>
<i>GNPOLD</i>	<i>VARCHAR2(9)</i>
<i>LOCALNAME</i>	<i>VARCHAR2(44)</i>
<i>GOVERNMENTFORM</i>	<i>VARCHAR2(44)</i>
<i>HEADOFSTATE</i>	<i>VARCHAR2(32)</i>
<i>CAPITAL</i>	<i>VARCHAR2(4)</i>
<i>CODE2</i>	<i>VARCHAR2(2)</i>

Solution Hint

Join the CITY and COUNTRY tables on the matching key columns (CITY.CountryCode and COUNTRY.Code) and sum the populations of cities where the continent is 'Asia'.

- Use a **JOIN** clause to combine CITY and COUNTRY tables based on matching key columns.
- Use a **WHERE** clause to filter for rows where CO.CONTINENT = 'Asia'.

SQL Query

```
1 SELECT SUM(ci.population) AS total_population
2 FROM CITY ci
3 JOIN COUNTRY co
4     ON ci.CountryCode = co.Code
5 WHERE co.Continent = 'Asia';
```

2.7 African Cities

Given the CITY and COUNTRY tables, query the names of all cities where the CONTINENT is 'Africa'.

Note: the input format and table structures are the same as the previous problem.

SQL Query

```
1 SELECT ci.name AS city_Afr
2 FROM city ci
3 JOIN country co
4     ON ci.countrycode = co.code
5 WHERE co.continent = 'Africa';
```

2.8 Average Population of Each Continent

Given the CITY and COUNTRY tables, query the names of all the continents (COUNTRY.Continent) and their respective average city populations (CITY.Population) rounded down to the nearest integer.

Note: the input format and table structures are the same as the previous problem.

Solution Hint

Using **GROUP BY clause**: is used to group rows that have **the same values** in specified columns into summary rows. It is often combined with aggregate functions like COUNT(), SUM(), AVG(), MAX(), and MIN() to produce summarized results./đề yêu cầu tính trung bình dân số theo từng châu lục, nên ta sẽ nhóm các thành phố theo châu lục rồi tính trung bình dân số của từng nhóm đó.

SQL Query

```
1 SELECT co.continent, FLOOR(AVG(ci.population))
2 FROM CITY ci
3 JOIN COUNTRY co
4     ON ci.CountryCode = co.Code
5 GROUP BY co.continent;
```

2.9 Draw The Triangle 1

P(R) represents a pattern drawn by Julia in R rows. The following pattern represents P(5):

```
* * * * *
* * * *
* * *
* *
*
```

Write a query to print the pattern P(20).

Solution Hint

Use a recursive common table expression (CTE) to generate numbers from 20 down to 1, and then use the REPEAT function to create the desired pattern for each number.

- WITH RECURSIVE nums AS (...): This defines a recursive CTE named nums that generates numbers starting from 20 down to 1.
- SELECT 20 AS n: This is the base case of the recursion, starting with the number 20.

- `UNION ALL SELECT n - 1 FROM nums WHERE n > 1`: This is the recursive part that decrements `n` by 1 until it reaches 1.
- `SELECT REPEAT('* ', n) AS pattern FROM nums`: This selects each number `n` from the CTE and uses the `REPEAT` function to generate a string of asterisks followed by spaces, repeated `n` times.

SQL Query

```

1 WITH RECURSIVE nums AS (
2     SELECT 20 AS n
3     UNION ALL
4     SELECT n - 1
5     FROM nums
6     WHERE n > 1
7 )
8 SELECT REPEAT('* ', n) AS pattern
9 FROM nums;
```

2.10 Draw The Triangle 2

`P(R)` represents a pattern drawn by Julia in `R` rows. The following pattern represents `P(5)`:

```

*
* *
* * *
* * * *
* * * * *
```

Write a query to print the pattern `P(20)`.

SQL Query

```

1 WITH RECURSIVE nums AS (
2     SELECT 1 AS n
3     UNION ALL
4     SELECT n + 1
5     FROM nums
6     WHERE n < 20)
7 SELECT REPEAT('* ', n) AS pattern
8 FROM nums;
```

3 Intermediate SQL

3.1 The PADS

Generate the following two result sets:

1. Query an alphabetically ordered list of all names in OCCUPATIONS, immediately followed by the first letter of each profession as a parenthetical (i.e.: enclosed in parentheses). For example: AnActorName(A), ADoctorName(D), AProfessorName(P), and ASingerName(S).
2. Query the number of occurrences of each occupation in OCCUPATIONS. Sort the occurrences in ascending order, and output them in the following format:

There are a total of [occupation_count] [occupation]s.

where [occupation_count] is the number of occurrences of an occupation in OCCUPATIONS and [occupation] is the lowercase occupation name. If more than one Occupation has the same [occupation_count], they should be ordered alphabetically.

Note: There will be at least two entries in the table for each type of occupation. **Input Format:** The OCCUPATIONS table is described as follows:

<i>Column</i>	<i>Type</i>
<i>name</i>	<i>String</i>
<i>occupation</i>	<i>String</i>

Occupations will only contain one of the following values: **Doctor, Professor, Singer, Actor**. **Sample Input:**

<i>name</i>	<i>occupation</i>
Samantha	Doctor
Julia	Actor
Maria	Actor
Meera	Singer
Ashely	Professor
Ketty	Professor
Christeen	Professor
Jane	Actor
Jenny	Doctor
Priya	Singer

Sample Output:

```

Ashely(P)
Christeen(P)
Jane(A)
Jenny(D)
Julia(A)
Ketty(P)
Maria(A)
Meera(S)
Priya(S)
Samantha(D)
There are a total of 2 doctors.
There are a total of 2 singers.
There are a total of 3 actors.
There are a total of 3 professors.

```

Explanation: The results of the first query are formatted to the problem description's specifications.

The results of the second query are ascendingly ordered first by number of names corresponding to each profession (2<=2<=3<=3), and then alphabetically by profession (Doctor <= Singer, and actor <= Professor).

Solution Hint:

First Query: Print names with the first letter of their occupation in parentheses, ordered alphabetically.

eg. Ashley + Professor => Ashley(P)

Syntax elements used:

- `CONCAT(string1, string2, ...)`: Concatenates multiple strings into one.
- `LEFT(string, number)`: Extracts the leftmost 'number' characters from 'string'.
- `ORDER BY column`: Sorts the result set based on the specified column.

SQL Query - First Query

```
1 SELECT CONCAT(Name, '(', LEFT(Occupation, 1), ')')
2 FROM OCCUPATIONS
3 ORDER BY Name;
```

Second Query: Count people in each occupation and format the output as specified.

Syntax elements used:

- `COUNT(*)`: Counts the number of rows in each group.
- `LOWER(string)`: Converts 'string' to lowercase.
- `GROUP BY column`: Groups rows that have the same values in the specified column.
- `ORDER BY column1, column2`: Sorts the result set first by 'column1', then by 'column2'.
- `CONCAT(string1, string2, ...)`: Concatenates multiple strings into one./ghép câu

SQL Query - Second Query

```
1 SELECT CONCAT('There are a total of ', COUNT(*), ' ', LOWER(
   Occupation), 's.')
2 FROM OCCUPATIONS
3 GROUP BY Occupation --group by occupation to count each occupation
4 ORDER BY COUNT(*), Occupation; --sort by count asc, then occupation
   alphabetically
```

3.2 Occupations

Pivot the Occupation column in OCCUPATIONS so that each Name is sorted alphabetically and displayed underneath its corresponding Occupation. The output should consist of four columns (Doctor, Professor, Singer, and Actor) in that specific order, with their respective names listed alphabetically under each column.

Note: Print NULL when there are no more names corresponding to an occupation.

Note: Input format and Sample Input are the same as the previous problem.

Sample Output:

Jenny	Ashley	Meera	Jane
Samantha	Christeen	Priya	Julia
NULL	Ketty	NULL	Maria

Explanation: The first column is an alphabetically ordered list of Doctor names.

The second column is an alphabetically ordered list of Professor names.

The third column is an alphabetically ordered list of Singer names.

The fourth column is an alphabetically ordered list of Actor names.

The empty cell data for columns with less than the maximum number of names per occupation (in this case, the Professor and Actor columns) are filled with NULL values.

Solution Hint

OCCUPATIONS table has 2 columns: Name and Occupation. We need to pivot the Occupation column to create separate columns for each occupation type (Doctor, Professor, Singer, Actor) and list the corresponding names under each column.

Syntax elements used:

- **PARTITION BY column:** Divides the output into partitions to which the ROW_NUMBER() function is applied./ở đây ta chia bảng thành 4 nhóm nghề khác nhau.
- **ROW_NUMBER():** Assigns a unique sequential integer to rows within a partition of a result set, starting at 1 for the first row in each partition./đánh số thứ tự cho từng tên trong mỗi nhóm nghề.

- ***CASE WHEN condition THEN result END***: Evaluates a list of conditions like an IF-THEN-ELSE statement.
- ***GROUP BY column***: Groups rows that have the same values in the specified column.
- ***ORDER BY column***: Sorts the result set based on the specified column.

SQL Query:

```

1 SELECT
2     MAX(CASE WHEN Occupation = 'Doctor' THEN Name END) AS
3         Doctor,
4     MAX(CASE WHEN Occupation = 'Professor' THEN Name END) AS
5         Professor,
6     MAX(CASE WHEN Occupation = 'Singer' THEN Name END) AS
7         Singer,
8     MAX(CASE WHEN Occupation = 'Actor' THEN Name END) AS
9         Actor
10 FROM (
11     SELECT
12         Name,
13         Occupation,
14         ROW_NUMBER() OVER (PARTITION BY Occupation ORDER
15             BY Name) AS rn
16     FROM OCCUPATIONS
17 ) ranked
18 GROUP BY rn
19 ORDER BY rn;

```

3.3 Binary Tree Nodes

You are given a table, BST, containing two columns: N and P, where N represents the value of a node in Binary Tree, and P is the parent of N.

<i>Column</i>	<i>Type</i>
<i>N</i>	<i>Integer</i>
<i>P</i>	<i>Integer</i>

Write a query to find the node type of Binary Tree ordered by the value of the node. Output one of the following for each node:

- **Root** If node is root node.
- **Leaf** If node is leaf node.
- **Inner** If node is neither root nor leaf node.


Sample Input:

N	P
1	2
3	2
6	8
9	8
2	5
8	5
5	NULL

Sample Output:

1 Leaf
 2 Inner
 3 Leaf
 5 Root
 6 Leaf
 8 Inner
 9 Leaf

Explanation: The Binary Tree below illustrates the sample:



Solution Hint

To determine the type of each node in the binary tree, we can use conditional statements to check the values of N and P.

- If P is NULL, the node is the root.
- If N does not appear as a parent (P) for any other node, it is a leaf.
- Otherwise, it is an inner node.
- Use ***CASE WHEN ... THEN ... ELSE ... END*** to implement the conditional logic.

SQL Query:

```
1 SELECT
2     N,
3     CASE
4         WHEN P IS NULL THEN 'Root'
5         WHEN N NOT IN (SELECT P FROM BST WHERE P IS NOT
6             NULL) THEN 'Leaf'
7         ELSE 'Inner'
8     END
9 FROM BST
ORDER BY N;
```

3.4 New Companies

Problem Statement: HackerRank - The Company

Solution Hint